


# Multiprocessing in Python

Tucson Python Meetup  
July 7th 2020

Ian Bertolacci  
[ian-bertolacci.github.io](https://ian-bertolacci.github.io)  
(Available for hire!)

Code and materials available at  
[github.com/ian-bertolacci/Multiprocessing-in-Python-Tucson-Python-Meetup](https://github.com/ian-bertolacci/Multiprocessing-in-Python-Tucson-Python-Meetup)



# Python Performance Disclaimer



- Python is not a high performance programming language.
  - Write in a faster language if you need extremely fast applications.
- Use fast libraries!
  - Have fast implementations (though unlikely to be parallel).
  - NumPy, SciPy, Pandas, etc

# Parallel Programming Disclaimer

- Parallelism is not always the answer
  - Amdahl's law:  $\text{speedup} \leq \text{total\_time} / ( \text{time\_in\_serial} + ( \text{time\_in\_parallel} / \text{processes} ) )$
  - If x% of a program is parallelizable, the **BEST CASE** is an x% reduction in runtime.
    - Period.
    - You will get less than that.
- More Parallelism != Faster Program
  - Overhead of spawning parallel tasks.
  - Resource contention.
- Profile! Experiment!
  - Is program slow because of amount of work? Or because implementation is algorithmically slow?



# Parallel Programming Disclaimer

"I SPEND A LOT OF TIME ON THIS TASK.  
I SHOULD WRITE A PROGRAM AUTOMATING IT!"



# Parallel Programming Disclaimer

"MY PROGRAM SPENDS A LOT OF TIME ON THIS TASK.  
I SHOULD PARALLELIZE IT!"



# Breaking Down a Problem into Parallelism

- What is "work"?
  - Multiple different tasks?
    - Implies task parallelism
  - Multiple identical tasks on different data?
    - Implies data parallelism
- Dependencies
  - What tasks depend on other each other?
    - Implies ordering between tasks
  - What tasks share objects?
    - Implies communication between tasks.
  - What tasks are modifying shared objects?
    - Implies synchronization between tasks

# Example: Adding Two Arrays

`added_AB = [ a + b for a, b in zip(A, B) ]`

- What kind of work is happening here?
  - Same task, different data.
- Are there dependencies between tasks?
  - No
  - All tasks operate independently.
- Do tasks share data?
  - No
  - All tasks have their own data.

# Example: Modifying, Sending, and Writing Data

```
fooized_data = fooize( data )  
send_data( fooized_data, server_url )  
json_write( fooized_data, "fooized.json" )
```

- What kind of work is happening here?
  - Different tasks
- Are there dependencies between tasks?
  - Yes
  - The `send_data` communication and `json_write` depend on the statement creating `fooized_data`
- Do tasks share data?
  - Yes
  - The server communication and file output share `fooized_data`
- Are tasks are modifying shared objects?
  - No
  - Because I made this example up and assume neither `send_data` or `json_write` modify the object



# Python Multiprocessing

- Standard Python library with API for parallel python programming.
  - Same API as Python Threading library (which is *\*not\** parallel).
- Excellent docs: [docs.python.org/3/library/multiprocessing.html](https://docs.python.org/3/library/multiprocessing.html)
- Provides:
  - Parallelism via Processes and Process pools
  - Synchronization via Locks, Events, Barriers, etc...
  - Inter-process communication via Pipes, Queues, Shared Memory

# Example: Data Parallelism using Pool.map

Original:

```
added_AB = [ a + b for a, b in zip(A, B) ]
```

Parallel:

with Pool( processes ) as pool:

```
    added_AB = pool.starmap( operator.add, zip(A, B) )
```

# Example: Parallel Reduction using Pool.map

Original:

```
sum_data = sum( data )
```

Parallel:

```
with Pool( processes ) as pool:
```

```
    chunk_size = int( ceil( len(data)/processes ) )
```

```
    chunked_work = [
```

```
        data[ i : i + chunk_size ]
```

```
        for i in range( 0, len(data), chunk_size )
```

```
    ]
```

```
    partial_sum_data = pool.imap( sum, chunked_work, chunksize=1 )
```

```
    sum_data = sum( partial_sum_data )
```

# Some Notes on Pool.map

- Careful about order!
  - if order matters, make sure you're using the correct map function
  - Especially true for reduction operations.
- A number of different map functions for different needs.
  - Asynchronous map
  - Starmap
  - Unordered map
  - Iterable map
  - Combinations of all
-

# Example: Task Parallelism with Pool.apply

Original:

```
fooized_data = fooize( data )  
send_data( fooized_data, server_url )  
json_write( fooized_data, "fooized.json" )
```

Parallel:

```
fooized_data = fooize( data )  
with Pool( processes ) as pool:  
    send_async = pool.apply_async( send_data, (fooized_data, server_url) )  
    write_async = pool.apply_async( json_write, (fooized_data, "fooized.json") )  
  
send_result = send_async.get()  
write_result = write_async.get()
```

# Example: Task Parallelism with Process

Original:

```
fooized_data = fooize( data )  
send_data( fooized_data, server_url )  
json_write( fooized_data, "fooized.json" )
```

Parallel:

```
fooized_data = fooize( data )  
processes = [  
    Process( target=send_data, args=(fooized_data, server_url) ),  
    Process( target=json_write, args=(fooized_data, "fooized.json") )  
]  
for process in processes: process.start()  
for process in processes: process.join()
```

# Process Communication and Shared Data

- Different tasks may be independent until they need to share data.
- Direct communication:
  - Pipes: connects two processes, can be bi-directional (Not covered here)
  - Queues: many-to-many FIFO queue
- Shared Data/Memory
  - Values: wraps a C type variable with a lock.
  - Arrays: wraps an array of C type values with a lock.

# Example: Queue

```
def make_work( queue, N ):  
    for work in range(N):  
        queue.put( work )  
    queue.put("dead")
```

```
def do_work( queue ):  
    while True:  
        work = queue.get()  
        if work == "dead":  
            break  
        print( f"got {work}" )
```

```
# main  
items = 100  
work_queue = Queue()  
Process( target = make_work, args =(work_queue, items) ).start()  
Process( target = do_work,    args =(work_queue, )      ).start()
```



# Example: Value

```
def writer( counter ):  
    with counter.get_lock():  
        while True:  
            counter.value += 1
```

```
def reader( counter ):  
    while True:  
        print( counter.value )
```

```
# main  
shared_value = Value( "i", 0 )  
Process( target = make_work, args =(shared_value, ) ).start()  
Process( target = do_work,    args =(shared_value, ) ).start()
```

# But What About Non-ctype Shared Memory?



# Synchronization

- Lock
  - Mutual exclusion (of code sections and/or data)
  - A process acquires the lock then releases the lock
- Event:
  - Many-on-one ordering
  - One-or-more processes wait on the event, a single process triggers it, unblocking all
- Condition:
  - Conditional mutual exclusion (think lock + event)
  - Process waits, another process unblocks
- Barriers:
  - Many-on-many ordering
  - All processes\* must reach and wait at the barrier before all can proceed
- Semaphore/BoundedSemaphore:
  - Capacity exclusion (think lock that can be opened by at most N keys simultaneously)
  - Only N processes may acquire the semaphore simultaneously

# Example: Lock

```
def worker( id, array ):  
    while True:  
        array[id] += 1
```

```
def printer( id, lock, array ):  
    while True:  
        sleep( 1 )  
        with lock:  
            for v in array:  
                print( f"({id}) {v}" )
```

```
# main  
array = Array( "i", [0] * workers )  
lock = Lock()  
[ Process( target = worker, args = (id, array) ).start() for id in ... ]  
[ Process( target = printer, args = (id, array, lock ) ).start() for id in ... ]
```

# Deadlocking

- Processes can acquire multiple locks.
- Deadlock occurs when two processes hold locks (a, b) while trying to acquire the other (b, a)
- Neither can progress, essentially halting the program.
- Easy to avoid if process will ever only hold one lock simultaneously.
- Hard to avoid otherwise.

```
def foo( lock_a, lock_b ):
    with lock_a:
        with lock_b:
            ... stuff ...
```

```
def bar( lock_a, lock_b ):
    with lock_b:
        with lock_a:
            ... stuff ...
```

# Example: Event and Barrier

```
def writer( id, array, barrier, event ):  
    while True:  
        event.wait()  
        while event.is_set():  
            array[id] += 1  
        barrier.wait()
```

```
# main  
event = mp.Event()  
write_barrier = mp.Barrier( writers +  
readers_barrier = mp.Barrier( readers )  
event.set()  
# ... the rest of the program ...
```

```
def safe_reader( array, wbarrier,  
rbarrier, event ):  
    while True:  
        sleep( 1 )  
        event.clear()  
        wbarrier.wait()  
        local_copy = [ v for v in array ]  
        rbarrier.wait()  
        event.set()
```

**Is this safe for multiple readers?**

**NO!**

# Parting Thoughts

- Is multiprocessing for everyone? Certainly not.
  - Still hard to use
  - Parallel engineering is weirder than in other languages (shared memory in particular)
- Is it useful? Mmmaybe?
  - Probably best for data parallelism
  - Hard to see where multiprocessing would be the best mechanism for task parallelism
    - Rapidly descends into microservices land
- Hints:
  - Read those docs!!!!
  - `Python.threading` and `Python.multiprocessing` have a very common API (intentional)
  - Be ready to kill some processes manually